

# Subquadratic Sparse Attention: Content-Routed Exact Attention for Long Contexts

Jon Smirl

June 21, 2026

## Abstract

Dense self-attention costs  $O(n^2)$  in the sequence length  $n$ , which is the binding constraint on long-context language models. We present *Subquadratic Sparse Attention* (SSA), an attention mechanism that, for each query, (i) routes to a small content-dependent set of key blocks using only per-block summary statistics, (ii) adds a local window, and (iii) performs *exact* softmax attention over the selected keys. The per-query work is  $O(\kappa)$  in a fixed budget  $\kappa \ll n$  plus a sublinear routing cost, so the layer runs in  $O(n\sqrt{n})$  flat or near-linear with a hierarchical router. We develop the theory that says when this is sound. A retrieval-margin analysis shows that softmax attention recovers a target key with weight  $\sigma(\beta\Delta - \log \mu)$ , where  $\Delta$  is the score margin and  $\mu$  the number of competing distractors; selection works because it cuts  $\mu$  from  $n$  to  $\kappa$ , making recovery *flat in  $n$* . We give the admissible routing bound that licenses summary-only selection, a tempered (cumulant) routing score that—unlike centroid routing—sees in-block outliers, and a closed-form variance prune test from Samuelson’s inequality. We then prove the limit: cheap, lossless, length-robust selection cannot hold simultaneously for arbitrary keys (a one-line probe argument), so SSA’s subquadratic *exactness* is licensed by the *benign geometry* of trained representations—geometry that training can be made to manufacture, via a routability regularizer that shrinks off-target spread. Finally we show length generalization (rotary position plus staged continued training reaches 32× the trained length at 0.98 recall for ~800 adaptation steps) and a construction pipeline that converts a dense pretrained model into a subquadratic one by swapping the attention and briefly adapting (recovering to within +1.2 perplexity of a dense model given equal training while attending 38% of keys). All claims are accompanied by measurements at controlled scale.

## 1 Introduction

A transformer layer computes, for queries  $Q \in \mathbb{R}^{n \times d}$ , keys  $K \in \mathbb{R}^{n \times d}$ , and values  $V \in \mathbb{R}^{n \times d}$ ,

$$\text{Attn}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right) V.$$

The  $n \times n$  score matrix makes both time and memory  $\Theta(n^2d)$ . For contexts of  $10^6$ – $10^7$  tokens this is prohibitive, yet most of the matrix is near-zero: for a given query only a small set of keys

carries appreciable weight. The question is whether one can *find* that set without forming all  $n^2$  scores.

Two subquadratic families answer differently. *Kernel / linear attention* replaces  $\exp(\langle q, k \rangle)$  by a factorizable feature map  $\phi(q)^\top \phi(k)$ , giving  $O(n)$  cost but a low-rank (smoothed) attention matrix. *Sparse / selective attention* keeps the exact softmax but evaluates it only on a chosen subset of keys. SSA is in the second family, with three design commitments:

1. **Content-dependent selection.** The chosen keys depend on the query, not only on position; this is what lets a query reach the one relevant block a million tokens back.
2. **Summary-only routing.** Selection uses per-block statistics (a mean, a covariance), so it does not read every key—that is where the subquadratic saving comes from.
3. **Exact attention over the selected set.** Within the selected keys the softmax is computed exactly, so there is no kernel-approximation error on the keys that matter.

The paper is organized around one tension. Selection is cheap only if it can judge a block from its summary; it is *correct* only if those summaries do not hide a key that mattered. Sections 3–5 make both sides precise; Section 6 shows they cannot both hold in the worst case, so something must give; Section 7 shows what gives—the data geometry, which training shapes. Sections 8–9 turn the mechanism into a usable recipe.

## 2 Notation and the retrieval view

Fix a single query  $q \in \mathbb{R}^d$  and keys  $k_1, \dots, k_n$ . Write the *logit* (score) of key  $j$  as

$$a_j = \langle q, k_j \rangle, \quad w_j = \frac{e^{\beta a_j}}{\sum_{l=1}^n e^{\beta a_l}}, \quad \beta = \frac{1}{\sqrt{d}},$$

so  $w$  is the attention distribution and  $\beta$  its inverse temperature. The output is  $o = \sum_j w_j v_j$ .

A long-context layer is, operationally, an *associative recall*: a query must place most of its weight on the key(s) that hold the information it needs and little on the rest. We therefore analyze attention by the weight it puts on a designated *target* key  $k_\star$  relative to the *distractors*  $\{k_j\}_{j \neq \star}$ .

## 3 Retrieval margin and the recovery weight

Define the *margin* of the target as the gap between its logit and the typical distractor logit,

$$\Delta = a_\star - \bar{a}_{\text{dist}}, \quad \bar{a}_{\text{dist}} = (\text{typical } a_j, j \neq \star).$$

Suppose there are  $\mu$  effective distractors with logit  $\approx a_\star - \Delta$ . Then the target weight is

$$w_\star = \frac{e^{\beta a_\star}}{e^{\beta a_\star} + \mu e^{\beta(a_\star - \Delta)}} = \frac{1}{1 + \mu e^{-\beta \Delta}} = \sigma(\beta \Delta - \log \mu), \quad (3.1)$$

where  $\sigma(x) = 1/(1 + e^{-x})$  is the logistic. Equation (3.1) is the *recovery weight*. Two consequences:

- **Detectability threshold.** The target is recovered ( $w_\star > \frac{1}{2}$ ) iff

$$\boxed{\beta \Delta > \log \mu} \quad (3.2)$$

Recovery degrades only *logarithmically* in the number of distractors.

- **Why selection helps, and why it is flat in  $n$ .** Dense attention pays  $\mu = n$ . A selector that attends only a budget of  $\kappa$  keys pays  $\mu = \kappa$ , replacing the threshold  $\log n$  by  $\log \kappa$ . If  $\kappa$  is held *fixed* as the context grows, the threshold (3.2) does not move with  $n$ : retrieval accuracy becomes *independent of context length*. This is the entire value proposition of selection, and it is why a fixed-budget selector can answer single-target queries at  $10^6$ – $10^7$  tokens that dense attention, with its  $\log n$  erosion, increasingly struggles with.

The catch is hidden in the phrase “a selector that attends  $\kappa$  keys”: the selector must *contain the target in its budget*. Sections 4–6 are about exactly that.

## 4 The algorithm

### 4.1 Block partition and summaries

Partition the  $n$  keys into  $B$  contiguous *blocks* of size  $b$  (so  $Bb = n$ ). For block  $c$  precompute, once, the summary statistics

$$\mu_c = \frac{1}{b} \sum_{j \in c} k_j \quad (\text{mean}), \quad \Sigma_c = \frac{1}{b} \sum_{j \in c} (k_j - \mu_c)(k_j - \mu_c)^\top \quad (\text{covariance}), \quad (4.1)$$

and a radius  $R_c = \max_{j \in c} \|k_j - \mu_c\|$ . In practice  $\Sigma_c$  is kept diagonal,  $\sigma_c^2 = \text{diag } \Sigma_c$ , so a block’s summary is  $O(d)$  numbers.

### 4.2 Routing score

For a query  $q$ , the *routing score* of block  $c$  estimates the largest logit the block can contribute. The exact quantity is the block’s log-sum-exp  $\text{LSE}_c(q) = \log \sum_{j \in c} e^{\beta \langle q, k_j \rangle}$ , which requires reading every key. SSA uses its *second-order (cumulant) surrogate*, computable from the summary alone:

$$r_c(q) = \langle q, \mu_c \rangle + \frac{\beta}{2} q^\top \Sigma_c q \approx \langle q, \mu_c \rangle + \frac{\beta}{2} \langle q^2, \sigma_c^2 \rangle. \quad (4.2)$$

Equation (4.2) is the Taylor/cumulant expansion of the tempered mean  $\beta^{-1} \log \frac{1}{b} \sum_{j \in c} e^{\beta \langle q, k_j \rangle}$  about  $\beta = 0$ : the first cumulant is the mean logit  $\langle q, \mu_c \rangle$ ; the second adds the *variance of the logit across the block*,  $q^\top \Sigma_c q$ . The variance term is what makes routing see an in-block outlier (Section 5.2).

### 4.3 Selection and exact attention

For each query  $q_i$  (at position  $i$ ):

1. Compute  $r_c(q_i)$  for all blocks  $c$  whose keys are causally visible ( $c$  ends at or before  $i$ ).
2. Select the top- $k$  blocks by  $r_c$ , *union* a local window of the  $w$  most recent blocks (recency is always relevant and cheap to include), giving a key set  $S_i$  with  $|S_i| = \kappa \leq (k + w)b$ .
3. Compute exact softmax attention restricted to  $S_i$ :

$$o_i = \sum_{j \in S_i} \frac{e^{\beta \langle q_i, k_j \rangle}}{\sum_{l \in S_i} e^{\beta \langle q_i, k_l \rangle}} v_j. \quad (4.3)$$

---

Algorithm 1: SSA forward (one head)

```

Input:  Q, K, V in  $\mathbb{R}^{n \times d}$ ;  block size  $b$ ;  budgets  $k$  (global),  $w$  (local)
Precompute: for each block  $c$ , mean  $\mu_c$  and diagonal covariance  $\Sigma_c$       //  $O(n d)$ 
For each query  $q_i$  (in parallel over  $i$ ):
  Routing:   for each causally visible block  $c$ :
               $r_c \leftarrow \langle q_i, \mu_c \rangle + (\beta/2) * \langle q_i^2, \text{diag } \Sigma_c \rangle$ 
  Selection:  $T \leftarrow (\text{top-}k \text{ blocks by } r_c) \cup (\text{w most-recent blocks})$ 
  Causal:     $S_i \leftarrow \text{keys of blocks in } T \text{ with position } \leq i$ 
  Attention:  $o_i \leftarrow \text{softmax}_{\{j \in S_i\}} ( \langle q_i, k_j \rangle / \sqrt{d} ) * V[S_i]$   // exact
Return 0

```

---

### 4.4 Complexity

Routing is  $O(n B d)$  if every query scores every block, and attention is  $O(n \kappa d)$ .

- **Flat router**,  $B = \sqrt{n}$ ,  $b = \sqrt{n}$ , fixed  $k, w$ : routing  $O(n^{1.5} d)$ , attention  $O(n \kappa d) = O(n^{1.5} d)$ . Total  $O(n^{1.5} d)$ —already a large saving over  $n^2$ .
- **Hierarchical router**. Organize blocks into a tree of summaries and descend it per query with branch-and-bound pruning (Section 5.1), so each query inspects  $O(k \log B)$  nodes rather than all  $B$ . Routing drops toward  $O(n \log n d)$  and the layer is near-linear in  $n$  up to the attention budget  $O(n \kappa d)$ .

Either way the dominant  $n^2$  term is gone. In a block-sparse kernel implementation the practical speedup over a dense exact kernel was  $20.6\times$  at  $n = 262,144$  on a single accelerator (Section 10).

## 5 Routing theory: when a summary is enough

The routing score (4.2) is a *heuristic* surrogate. This section gives the *certified* objects—upper bounds that make selection provably lossless—and the prune test SSA uses.

### 5.1 The admissible bound and branch-and-bound exactness

For any key  $k$  in block  $c$ , decompose its logit around the block mean and apply Cauchy–Schwarz:

$$\langle q, k \rangle = \langle q, \mu_c \rangle + \langle q, k - \mu_c \rangle \leq \langle q, \mu_c \rangle + \|q\| \|k - \mu_c\| \leq \underbrace{\langle q, \mu_c \rangle + \|q\| R_c}_{U_c(q)}. \quad (5.1)$$

$U_c(q)$  is an *admissible upper bound*: no key in block  $c$  has a logit exceeding it, and it depends only on the summary  $(\mu_c, R_c)$ . This licenses exact selection at adaptive cost.

**Branch-and-bound selection.** Maintain the best *actual* logit found so far,  $s^*$ . Open blocks in decreasing  $U_c(q)$ . Stop as soon as the next block has  $U_c(q) \leq s^*$ : by (5.1) no unopened block can contain a key beating  $s^*$ . The result is identical to scanning all keys, but only the blocks whose bound clears  $s^*$  are ever read.

The cost of branch-and-bound is the number of blocks whose bound exceeds the true best—a quantity governed entirely by how *tight* the bound (5.1) is, i.e. by the geometry of the keys (Sections 6–7).

**Anisotropic refinement.** The isotropic radius  $R_c$  is loose when a block’s keys are spread unevenly across directions. Using the covariance ellipsoid instead,

$$\langle q, k - \mu_c \rangle \leq \sqrt{q^\top \Sigma_c q} \cdot \rho_c, \quad \rho_c = \max_{j \in c} \|\Sigma_c^{-1/2}(k_j - \mu_c)\|, \quad (5.2)$$

which replaces  $\|q\|R_c$  by the *directional radius*  $\sqrt{q^\top \Sigma_c q} \cdot \rho_c$ . When the keys are anisotropic (the usual trained case),  $\sqrt{q^\top \Sigma_c q}$  can be far smaller than  $\|q\|R_c$  for queries pointing along thin directions—a tighter, query-specific bound, and exactly the quantity the cumulant score (4.2) already trades on.

### 5.2 Why second order: centroid routing is blind to outliers

Centroid routing uses only  $r_c = \langle q, \mu_c \rangle$ . A single key  $k_\star$  in a block of size  $b$  contributes  $\frac{1}{b}k_\star$  to  $\mu_c$ , so its signal in the mean is attenuated by  $1/b$ : a lone target in a large block is invisible

to centroid routing. The variance term in (4.2) repairs this, because an outlier aligned with  $q$  inflates  $q^\top \Sigma_c q$ . Concretely, if block  $c$  holds one key with  $\langle q, k_\star \rangle = a_\star$  and  $b - 1$  keys with logit  $\approx 0$ , then  $\langle q, \mu_c \rangle \approx a_\star/b$  but  $q^\top \Sigma_c q \approx a_\star^2/b$ , so the second-order score  $r_c \approx a_\star/b + (\beta/2)a_\star^2/b$  carries the quadratic outlier signal the mean discards.

**The tempered family and its bias.** Routing by the tempered score  $r_c^{(\beta)} = \beta^{-1} \log \sum_{j \in c} e^{\beta \langle q, k_j \rangle}$  interpolates between the mean ( $\beta \rightarrow 0$ ) and the max ( $\beta \rightarrow \infty$ ), and is sandwiched by

$$\max_{j \in c} \langle q, k_j \rangle \leq r_c^{(\beta)} \leq \max_{j \in c} \langle q, k_j \rangle + \frac{\log b}{\beta}. \quad (5.3)$$

So  $r_c^{(\beta)}$  estimates the block’s *best* logit—exactly what selection wants—with bias at most  $(\log b)/\beta$ . Small  $\beta$  smooths over outliers (the centroid failure); large  $\beta$  removes the bias but amplifies noise and loses the averaging that makes summaries stable. The cumulant form (4.2) is the second-order member of this family; empirically the routing quality is maximized near  $\beta \approx 2$ .

### 5.3 The Samuelson prune test

A closed-form, summary-only test for *discarding* a block uses *Samuelson’s inequality*: for any reals  $s_1, \dots, s_m$  with mean  $\bar{s}$  and population variance  $\text{Var} = \frac{1}{m} \sum_j (s_j - \bar{s})^2$ , every element obeys

$$(s_i - \bar{s})^2 \leq (m - 1) \text{Var}, \quad \text{equivalently} \quad \max_j s_j \leq \bar{s} + \sqrt{(m - 1) \text{Var}}. \quad (5.4)$$

Apply it to the in-block logits  $s_j = \langle q, k_j \rangle$ , whose mean is  $\langle q, \mu_c \rangle$  and whose variance is  $q^\top \Sigma_c q$ . Then the block’s best logit is bounded by  $\langle q, \mu_c \rangle + \sqrt{(b - 1) q^\top \Sigma_c q}$ , giving the *prune gate*: block  $c$  can be safely discarded against a threshold  $\tau = s^\star$  whenever

$$\boxed{(s^\star - \langle q, \mu_c \rangle)^2 > (b - 1) q^\top \Sigma_c q \quad \text{and} \quad \langle q, \mu_c \rangle < s^\star} \quad (5.5)$$

i.e. when the *margin* of the current best over the block mean exceeds  $\sqrt{(b - 1) \cdot \text{spread}}$ . The test needs only  $(\mu_c, \Sigma_c)$ . It is *sufficient* (it never wrongly prunes) but not necessary, and it sharpens the radius bound (5.1) by using the variance rather than the worst-case radius. Equation (5.5) is the operational core of cheap exact selection: it fires—and the block is skipped—precisely when the off-target spread  $q^\top \Sigma_c q$  is small, which is the benign-geometry condition of Section 7.

## 6 The trilemma and the impossibility

Call a selector *cheap* if it reads  $o(n)$  keys, *lossless* if it attends every key dense attention would weight non-negligibly, and *length-robust* if its accuracy is flat in  $n$ . The bounds above let us state the fundamental limit.

**Proposition 1** (No free selection). *No selector can be simultaneously cheap and lossless for arbitrary keys.*

*Proof.* Suppose a selector reads a set  $\mathcal{R}$  of keys with  $|\mathcal{R}| < n$ , and let  $j_0 \notin \mathcal{R}$ . Construct a probe input identical on  $\mathcal{R}$  but with  $k_{j_0} = cq$  for  $c$  large. Dense attention puts weight  $\rightarrow 1$  on  $j_0$ , so the correct output is  $v_{j_0}$ . The selector, never having read  $j_0$ , returns the same output as on the unmodified input, which is independent of  $v_{j_0}$ . Hence it is lossy on this input.  $\square$

Proposition 1 says losslessness for *worst-case* keys forces  $|\mathcal{R}| = n$ —no summary suffices, because a summary can always hide a spike. A quantitative companion holds under fine-grained complexity assumptions: exact attention with unbounded logits requires  $n^{2-o(1)}$  time, and only bounded-logit / low-rank regimes admit truly subquadratic exact computation. So at most *two* of {cheap, lossless, length-robust} are available at once:

| keep                     | give up       | what it is                                                    |
|--------------------------|---------------|---------------------------------------------------------------|
| lossless + length-robust | cheap         | dense attention—reads every key                               |
| cheap + length-robust    | lossless      | approximate selection—fast, flat, can miss a worst-case spike |
| cheap + lossless         | length-robust | branch-and-bound on <i>benign</i> keys—bounds tight, low cost |

The escape from the trilemma is the third row: cheapness *and* losslessness are jointly available *when the bounds* (5.1)/(5.2)/(5.5) *are tight*, which is a property of the key geometry, not of the algorithm. SSA is therefore not a universal subquadratic exact attention—no such thing exists—but a mechanism that is cheap, lossless, *and* length-robust *on benign geometry*, and merely cheap-and-length-robust-but-approximate otherwise. The next section is about making the geometry benign.

## 7 Manufacturing routability

### 7.1 Benign geometry, precisely

The branch-and-bound cost and the prune gate (5.5) are governed by the *off-target spread*  $q^\top \Sigma_c q$  for the blocks a query does *not* need, relative to the *margin*  $\Delta$  to the block it does. Geometry is *benign* for a query  $q$  when, for every non-target block  $c$ ,

$$\langle q, \mu_c \rangle + \sqrt{(b-1) q^\top \Sigma_c q} < s^*(q), \quad (7.1)$$

i.e. the prune gate (5.5) fires everywhere except the target’s block. Then exactly one block (plus the local window) is opened, branch-and-bound reads  $O(b)$  keys, and selection is cheap *and* lossless. Isotropic keys violate (7.1)—every block’s bound is large and uninformative, so nothing prunes and cost reverts to dense.

## 7.2 The routability regularizer

Benign geometry is not given; it can be *trained in*. Add to the language-model loss a penalty that shrinks the off-target spread along the directions queries actually use:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \lambda \sum_i \sum_{c \neq \text{target}(i)} q_i^\top \Sigma_c q_i. \quad (7.2)$$

Minimizing  $\sum_{c \neq \text{target}} q^\top \Sigma_c q$  is exactly minimizing the quantity that appears under the square root in the prune gate (5.5), so as training proceeds the gate fires for more non-target blocks and the certified bound (5.1)/(5.2) tightens. Measured: co-training with (7.2) drove the *lossless* branch-and-bound cost from 26.5% of keys to 4.2%—a  $6\times$  reduction—at *no* loss of retrieval accuracy. Training is what manufactures the geometry that the impossibility result (Proposition 1) says cheap exactness requires.

## 7.3 The entry-magnitude split: selection vs. linear attention

A complementary fact decides *which* subquadratic route a given layer should take. Let  $B$  be the characteristic magnitude of the logits  $\langle q, k \rangle$  (the “entry scale”). The exponentiated score matrix  $[e^{B\langle q_i, k_j \rangle}]$  has effective rank that *grows with*  $B$ : for small  $B$  it is nearly low-rank (so a kernel/linear-attention factorization  $\phi(q)^\top \phi(k)$  is accurate and  $O(n)$ ); for large  $B$  it is full-rank (so no linear map approximates it, and one must *select*). Empirically, the effective rank of  $e^{BKK^\top}$  on 256 keys rose from 2.7 at  $B = 0.5$  to 256 (full) by  $B = 8$ .

The selection route, by contrast, is *scale-invariant*: in the prune gate (5.5) both the squared margin and the spread scale as  $B^2$ , so the gate’s truth value is unchanged by  $B$ . Selection therefore works at *any* entry scale, whereas linear attention works only in the small- $B$  (smooth) regime. Sharp, long-range retrieval is the large- $B$  regime—which is why selection, not linearization, is the route for long-context exact recall.

# 8 Length generalization and staged extension

A selector that is flat in  $n$  (Section 3) still needs the *model* to behave at lengths beyond those it was trained on. The position encoding decides whether it can.

## 8.1 Position-invariant routing under rotary embeddings

With rotary position embeddings (RoPE), a query/key pair interacts only through their *relative* offset: the logit  $\langle q_i, k_j \rangle$  depends on  $i - j$ , not on  $i, j$  absolutely. A model that has learned *content* routing—match a query to the key whose content binds it, at whatever offset—therefore transfers to offsets it never saw, because (i) the decisive relative structure (e.g. a key-to-value offset of  $+1$ ) is constant at any length and (ii) content matching is position-free. A model with *learned*



*absolute* position embeddings cannot: its embeddings past the trained length are untrained. This predicts, and we observe, zero-shot extrapolation of  $\sim 2\text{--}4\times$  for RoPE and immediate collapse for learned-absolute position.

## 8.2 The staging ladder

Zero-shot extrapolation buys a factor, not an unbounded range. To go further, *stage*: extend the context, briefly continue-train (adapt) at the new length, extend again—doubling each rung. The economics are favorable precisely because routing is position-invariant: each rung extrapolates only  $2\times$  from the last *adapted* length, so every adaptation starts from the  $\sim 2\times$  zero-shot recall and only has to clean it up. The adaptation cost is small and roughly *flat in length* rather than growing.

Measured (a 6-layer model on a synthetic multi-query associative-recall task; base trained at 48 pairs):

| rung | multiple   | zero-shot recall | adapt steps | recall after adapt | recall with SSA |
|------|------------|------------------|-------------|--------------------|-----------------|
| 96   | $2\times$  | 0.993            | 100         | 1.000              | 1.000           |
| 192  | $4\times$  | 0.904            | 100         | 0.999              | 1.000           |
| 384  | $8\times$  | 0.810            | 100         | 1.000              | 0.995           |
| 768  | $16\times$ | 0.672            | 100         | 0.996              | 0.996           |
| 1536 | $32\times$ | 0.516            | 400         | 0.982              | 0.979           |

The base schedule cost 7600 steps; the *entire* climb from  $16\times$  (where zero-shot recall is already near the floor) to  $32\times$  at recall 0.982 cost only 800 additional steps—about 10% of the base, and roughly constant per rung. The final column confirms that the same adapted model runs under SSA selection ( $\leq 15\%$  of keys attended) with essentially no recall loss at every rung.

## 9 The construction pipeline: dense to subquadratic

SSA can be retrofitted onto an existing dense pretrained model, which is how a frontier long-context system is built without pretraining from scratch:

1. **Start** from a dense pretrained base.
2. **Swap** dense attention for SSA (Algorithm 1) in every layer.
3. The swap **degrades** quality, because the base was trained expecting every key.
4. **Adapt** by continued pretraining; the keys co-adapt to the sparse routing (and, with the regularizer (7.2), become routable).
5. **Stage** the context extension (Section 8).

Measured on a 124M-parameter dense base, held-out perplexity (lower is better):

| condition                                                 | perplexity  |
|-----------------------------------------------------------|-------------|
| dense base (off-domain start)                             | 32.5        |
| + SSA swap, no adaptation                                 | 45.8        |
| dense base + equal in-domain continued training (control) | 23.9        |
| <b>SSA-swapped + equal in-domain continued training</b>   | <b>25.1</b> |

The swap costs +13.2 perplexity; an *equal-budget* adaptation recovers 94% of that gap, landing within +1.2 of a dense model given the *same* continued training—while attending only  $\sim 38\%$  of keys at this configuration. The control matters: continued training lowers perplexity for either attention (domain adaptation), so “recovery” must be measured against a dense model given the same steps, not against the off-domain start. The residual is the price of sparsity at this small budget and shrinks as the budget grows.

## 10 Experiments

All experiments are at controlled scale; the point is to validate mechanisms, not to set absolute records.

**Routing.** Second-order (cumulant) routing recovers targets where centroid routing collapses, matching the  $1/b$  outlier-attenuation analysis of Section 5.2; routing quality peaks near  $\beta \approx 2$ , consistent with the bias–variance reading of (5.3).

**Kernel speedup.** A block-sparse implementation of Algorithm 1 achieved a  $20.6\times$  wall-clock speedup over a dense exact kernel at  $n = 262,144$  on a single accelerator, with the crossover well below that length.

**Routability.** The regularizer (7.2) reduced lossless branch-and-bound selection cost from 26.5% to 4.2% of keys at zero accuracy cost. The reduction was robust across head dimension and showed no capacity trade-off down to  $d =$  (number of clusters), since query-specific anisotropy needs only  $\sim 1$  dimension per cluster—real head dimensions (64–256) sit well above this.

**Length generalization and staging.** See Section 8:  $32\times$  the trained length at recall 0.982 for  $\sim 800$  adaptation steps, SSA-preserved.

**Construction pipeline.** See Section 9: swap +13.2 perplexity, 94% recovered to within +1.2 of the dense-adapted control at  $\sim 38\%$  of keys.

**What the headline retrieval numbers do and do not show.** Reported single-target needle-in-a-haystack accuracies of 98–100% at  $10^6$ – $10^7$  tokens are real and consistent with (3.1)–(3.2), but they live in a specific regime. We measured retrieval as a function of context length at fixed budget  $\kappa \approx 10^3$  and margin  $\Delta = 0.55$ :

| context $n$ | dense | SSA, isolated target | SSA, benign target |
|-------------|-------|----------------------|--------------------|
| 1024        | 1.00  | 1.00                 | 1.00               |
| 4096        | 1.00  | 0.47                 | 1.00               |
| 16384       | 1.00  | 0.20                 | 1.00               |
| 65536       | 0.90  | 0.10                 | 1.00               |
| 262144      | 0.83  | 0.00                 | 1.00               |

An *isolated* target—a lone spike with no correlated neighbors—*collapses* with length: cheap moment routing averages it into its block and loses it among the fluctuations of the growing number of blocks (the  $1/b$  attenuation of Section 5.2, now competing against more and more random blocks). This is Proposition 1 in miniature. A *benign* target—one accompanied by a coherent span of query-aligned neighbors, as a real answer is by its surrounding context—lifts its whole block’s score and stays flat at 1.00, *beating dense* at long  $n$  because selection caps the effective distractor count. The same separation appears across margins: an isolated target barely fires at any margin, while a benign one tracks dense once the margin clears the budget floor  $\sqrt{2 \log \kappa / d}$ . So the headline numbers certify the *easy and benign* regime (single, high-margin, accuracy not losslessness)—the regime everyone agrees is achievable—and do not certify lossless selection in the worst case, which Proposition 1 shows is unavailable cheaply.

## 11 Discussion and limitations

SSA is best understood as the resolution of a constrained problem rather than a universal accelerator. The recovery-weight law (3.1) shows selection makes retrieval flat in  $n$ ; the admissible bound (5.1) and the prune gate (5.5) show summaries can certify that selection losslessly; the trilemma (Section 6) shows this can be cheap *only* on benign geometry; and the regularizer (7.2) shows training can supply that geometry. The construction pipeline (Section 9) and the staging ladder (Section 8) turn the mechanism into a recipe that retrofits a dense model and extends its context a rung at a time.

Honest limitations follow directly from the theory. (i) *Worst-case losslessness is not available cheaply* —a sufficiently adversarial or genuinely low-margin target can always evade summary routing; SSA’s guarantees are conditional on the (trained, measured) benign geometry. (ii) *Multi-needle and low-margin retrieval* are the hard regime the headline single-needle numbers do not address. (iii) The selection budget  $\kappa$  sets a floor margin  $\sqrt{2 \log \kappa / d}$ ; targets below it are missed regardless of  $n$ . (iv) The experiments here are at modest scale ( $10^2$ – $10^5$  tokens,  $10^8$ -parameter

base); they validate mechanisms, and reaching  $10^7$ -token contexts is the same construction repeated with the hierarchical router, more compute, and a long-context training corpus. None of these is a missing algorithmic ingredient; they are the boundary the theory itself draws.

## A Derivations

**A.1 Recovery weight (3.1).** With one target at logit  $a_\star$  and  $\mu$  distractors at  $a_\star - \Delta$ ,  $w_\star = e^{\beta a_\star} / (e^{\beta a_\star} + \mu e^{\beta(a_\star - \Delta)})$ . Divide numerator and denominator by  $e^{\beta a_\star}$  to get  $1/(1 + \mu e^{-\beta \Delta}) = \sigma(\beta \Delta - \log \mu)$ , since  $1/(1 + e^{-x}) = \sigma(x)$  with  $x = \beta \Delta - \log \mu$ .

**A.2 Cumulant score (4.2).** Let  $g(\beta) = \beta^{-1} \log \frac{1}{b} \sum_{j \in c} e^{\beta \langle q, k_j \rangle}$ . As  $\beta \rightarrow 0$ ,  $g(\beta) = \mathbb{E}_j \langle q, k_j \rangle + \frac{\beta}{2} \text{Var}_j \langle q, k_j \rangle + O(\beta^2)$ , the cumulant generating function expansion. The mean is  $\langle q, \mu_c \rangle$  and the variance is  $q^\top \Sigma_c q$ , giving (4.2).

**A.3 Log-sum-exp sandwich (5.3).** For any reals  $x_1, \dots, x_b$  with  $M = \max_j x_j$ :  $e^{\beta M} \leq \sum_j e^{\beta x_j} \leq b e^{\beta M}$ . Take log, divide by  $\beta$ :  $M \leq \beta^{-1} \log \sum_j e^{\beta x_j} \leq M + \beta^{-1} \log b$ .

**A.4 Samuelson bound (5.4).** Center the data,  $d_j = s_j - \bar{s}$ , so  $\sum_j d_j = 0$ . Fix index  $i$ . By Cauchy–Schwarz over the other  $m - 1$  indices,  $d_i^2 = (\sum_{j \neq i} d_j)^2 \leq (m - 1) \sum_{j \neq i} d_j^2 = (m - 1)(\sum_j d_j^2 - d_i^2)$ . Hence  $d_i^2 m \leq (m - 1) \sum_j d_j^2$ , i.e.  $(s_i - \bar{s})^2 \leq (m - 1) \text{Var}$ . Taking the max over  $i$  and adding  $\bar{s}$  gives the stated bound on  $\max_j s_j$ , and (5.5) is its contrapositive against the threshold  $s^\star$ .

## References

- [1] A. Vaswani et al. *Attention Is All You Need*. NeurIPS, 2017.
- [2] R. Child, S. Gray, A. Radford, I. Sutskever. *Generating Long Sequences with Sparse Transformers*. 2019.
- [3] I. Beltagy, M. E. Peters, A. Cohan. *Longformer: The Long-Document Transformer*. 2020.
- [4] M. Zaheer et al. *Big Bird: Transformers for Longer Sequences*. NeurIPS, 2020.
- [5] N. Kitaev, Ł. Kaiser, A. Levskaya. *Reformer: The Efficient Transformer*. ICLR, 2020.
- [6] A. Roy, M. Saffar, A. Vaswani, D. Grangier. *Efficient Content-Based Sparse Attention with Routing Transformers*. TACL, 2021.
- [7] A. Katharopoulos, A. Vyas, N. Pappas, F. Fleuret. *Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention*. ICML, 2020.

- [8] K. Choromanski et al. *Rethinking Attention with Performers*. ICLR, 2021.
- [9] J. Su et al. *RoFormer: Enhanced Transformer with Rotary Position Embedding*. 2021.
- [10] L. Greengard, V. Rokhlin. *A Fast Algorithm for Particle Simulations*. J. Comput. Phys., 1987.
- [11] J. Barnes, P. Hut. *A Hierarchical  $O(N \log N)$  Force-Calculation Algorithm*. Nature, 1986.
- [12] P. A. Samuelson. *How Deviant Can You Be?* J. Amer. Statist. Assoc., 1968.
- [13] A. Keles, P. Wijewardena, C. Hegde. *On the Computational Complexity of Self-Attention*. ALT, 2023.
- [14] J. Alman, Z. Song. *Fast Attention Requires Bounded Entries*. NeurIPS, 2023.
- [15] H. Ramsauer et al. *Hopfield Networks Is All You Need*. ICLR, 2021.
- [16] S. Arora et al. *Zoology: Measuring and Improving Recall in Efficient Language Models*. 2023.